

# LAB 4

## AMIL

Student Name: Joud Alahmari

Student ID: 2008071

Import data set

And do the splitting of data set and training set

Note: 20% for testset, 80% for training set

Spyder (Python 3.9)

File Edit Search Source Run Debug Consoles Projects Tools View Help

C:\Users\joudm\Desktop\ML\LAB4\_2008071\_spyder.py

LAB4\_2008071\_spyder.py X lab3\_ML\_2008071\_joud.py X

```

1  #-*- coding: utf-8 -*-
2  """
3  Created on Mon Jan  9 12:40:58 2023
4
5  @author: joud Alahmari
6  """
7
8  # 1-import Libraries
9  import numpy as np
10 import matplotlib.pyplot as plt
11 import pandas as pd
12
13 # 2-import dataSet
14 mydataset= pd.read_csv("C:\\Users\\joudm\\Desktop\\ML\\Social_Network_Ads.csv")
15 X = mydataset.iloc[:,[2,3]].values
16 y = mydataset.iloc[:, -1].values
17

```

mydataset - DataFrame

Index	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15683246	Female	27	57000	0
4	15884002	Male	19	76000	0
5	15728773	Male	27	58000	0
6	15598044	Female	27	84000	0
7	15694829	Female	32	150000	1
8	15600575	Male	25	33000	0
9	15727311	Female	35	65000	0
10	15570769	Female	26	80000	0
11	15666374	Female	26	52000	0

Variable Explorer

Python Console

History

X - NumPy object array

	0	1
0	19	19000
1	35	20000
2	26	43000
3	27	57000
4	19	76000
5	27	58000
6	27	84000
7	32	150000
8	25	33000
9	35	65000
10	26	80000
11	26	52000

y - NumPy object array

	0
0	0
1	0
2	0
3	0
4	0
5	0
6	0
7	1
8	0
9	0
10	0
11	0

Training the Logistic Regression model on the Trainingset  
And  
predict the results

# MACHINE LEARNING

Spyder (Python 3.9)

File Edit Search Source Run Debug Consoles Projects Tools View Help

C:\Users\joumd\Desktop\ML\LAB4\_2008071\_spyder.py

```
LAB4_2008071_spyder.py x lab3_ML_2008071_joud.py x
8 # 1-import Libraries
9 import numpy as np
10 import matplotlib.pyplot as plt
11 import pandas as pd
12
13 # 2-import DataSet
14 mydataset= pd.read_csv("C:\\Users\\joumd\\Desktop\\ML\\Social_Network_Ads.csv")
15 X = mydataset.iloc[:,2:3].values
16 y = mydataset.iloc[:, -1].values
17
18 # 4-splitting the dataset into training set and test set
19 from sklearn.model_selection import train_test_split
20 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.20, random_state = 0 )
21
22 # 5-feature Scaling
23 from sklearn.preprocessing import StandardScaler
24 sc = StandardScaler()
25 X_train = sc.fit_transform(X_train)
26 X_test = sc.fit_transform(X_test)
27
28 # 6-Training the Logistic Regression model on the Training set
29
30 from sklearn.linear_model import LogisticRegression
31
32 classifier = LogisticRegression(random_state= 0)
33
34 classifier.fit(X_train, y_train)
35
36
37
38
39 # 7-Predicting the Test set results
40
41 y_pred = classifier.predict(X_test)
42
43
44
45
46
47
```

Name	Type	Size	Value
classifier	linear_model.logistic.LogisticRegression	1	LogisticRegression object of sklearn.linear_model.logistic module
mydataset	DataFrame	(400, 5)	Column names: User ID, Gender, Age, EstimatedSalary, Purchased
sc	preprocessing_data.StandardScaler	1	StandardScaler object of sklearn.preprocessing_data module
X	Array of int64	(400, 2)	[[ 19 19000] [ 35 20000]
X_test	Array of float64	(80, 2)	[[ -0.49618606  0.56021375] [ 0.2389044  -0.59133674]
X_train	Array of float64	(320, 2)	[[ 1.92295088  2.14601566] [ 2.02016082  0.3787193 ]]
y	Array of int64	(400,)	[0 0 0 ... 1 0 1]
y_pred	Array of int64	(80,)	[0 0 0 ... 0 1 1]
y_test	Array of int64	(80,)	[0 0 0 ... 0 1 1]
y_train	Array of int64	(320,)	[1 0 0 ... 0 0 0]

Help Variable Explorer Plots Files

Console 1/A x

Python 3.9.13 (main, Aug 25 2022, 23:51:50) [MSC v.1916 64 bit (AMD64)]  
Type "copyright", "credits" or "license" for more information.

IPython 7.31.1 -- An enhanced Interactive Python.

In [1]: runfile('C:/Users/joumd/Desktop/ML/LAB4\_2008071\_spyder.py', wdir='C:/Users/joumd/Desktop/ML')

In [2]: runfile('C:/Users/joumd/Desktop/ML/LAB4\_2008071\_spyder.py', wdir='C:/Users/joumd/Desktop/ML')

In [3]:

Python Console History

LSP Python: ready conda: base (Python 3.9.13) Line 47, Col 1 UTF-8 CRLF RW Mem 67%

y\_train - NumPy object array

	0
0	1
1	0
2	0
3	0
4	1
5	1
6	0
7	1
8	1
9	0
10	0
11	1

y\_test - NumPy object array

	0
0	0
1	0
2	0
3	0
4	0
5	0
6	0
7	1
8	0
9	0
10	0
11	0

X\_test - NumPy object array

	0	1
0	-0.496186	0.560214
1	0.238904	-0.591337
2	-0.0367545	0.186738
3	-0.496186	0.31123
4	-0.0367545	-0.591337
5	-0.771845	-1.52503
6	-0.4043	-1.68064
7	0.0551318	2.33422
8	-1.59882	-0.031123
9	1.06588	-0.809198
10	-0.496186	-0.62246
11	-0.679959	-0.435722

X\_train - NumPy object array

	0	1
0	1.92295	2.14602
1	2.02016	0.378719
2	-1.38222	-0.432499
3	-1.18779	-1.01194
4	1.92295	-0.925024
5	0.367578	0.291803
6	0.173157	0.146943
7	2.02016	1.74041
8	0.756421	-0.838108
9	0.270367	-0.287638
10	0.367578	-0.17175
11	-0.118476	2.20396

y\_pred - NumPy object array

	0
0	0
1	0
2	0
3	0
4	0
5	0
6	0
7	1
8	0
9	1
10	0
11	0

Format Resize Background color

## Visualising the Training set results

Spyder (Python 3.9)

File Edit Search Source Run Debug Consoles Projects Tools View Help

C:\Users\youdm\Desktop\ML

C:\Users\youdm\Desktop\ML\LAB4\_2008071\_spyder.py

```

26 X_test = sc.fit_transform(X_test)
27
28
29 # 6-Training the Logistic Regression model on the Training set
30 from sklearn.linear_model import LogisticRegression
31 classifier = LogisticRegression(random_state=0)
32 classifier.fit(X_train, y_train)
33
34
35 # 7-Predicting the Test set results
36 y_pred = classifier.predict(X_test)
37
38
39 # 8-Making the Confusion Matrix
40 from sklearn.metrics import confusion_matrix
41 cm = confusion_matrix(y_test, y_pred)
42 print(cm)
43
44
45 # 9-Visualising the Training set results
46 from matplotlib.colors import ListedColormap
47 X_set, y_set = X_train, y_train
48 X1, X2 = np.meshgrid(np.arange(start = X_set[:, 0].min() - 1, stop = X_set[:, 0].max() + 1, step = 0.01),
49 np.arange(start = X_set[:, 1].min() - 1, stop = X_set[:, 1].max() + 1, step = 0.01))
50 plt.contourf(X1, X2, classifier.predict(np.array([X1.ravel(), X2.ravel()]).T).reshape(X1.shape),
51 alpha = 0.75, cmap = ListedColormap(('red', 'green')))
52 plt.xlim(X1.min(), X1.max())
53 plt.ylim(X2.min(), X2.max())
54 for i, j in enumerate(np.unique(y_set)):
55     plt.scatter(X_set[y_set == j, 0], X_set[y_set == j, 1], alpha=0.75,
56 c = ListedColormap(('red', 'green'))(i), label = j)
57 plt.title('Logistic Regression (Training set)')
58 plt.xlabel('Age')
59 plt.ylabel('Estimated Salary')
60 plt.legend()
61 plt.show()
62
63
64
65

```

Name	Type	Size	Value
classifier	linear_model._logistic.LogisticRegression	1	LogisticRegression object of sklearn
cm	Array of int64	(2, 2)	[[52 6] [ 3 19]]
i	int	1	1
j	int64	1	1
mydataset	DataFrame	(400, 5)	Column names: User ID, Gender, Age, EstimatedSalary, Purchased
sc	preprocessing._data.StandardScaler	1	StandardScaler object of sklearn.preprocessing._data module
X	Array of int64	(400, 2)	[[ 19 19000] [ 35 20000]]
X1	Array of float64	(592, 609)	[[ -2.96547978 -2.95547978 -2.94547...
X2	Array of float64	(592, 609)	[[ -2.59138156 -2.59138156 -2.59138...
X_set	Array of float64	(320, 2)	[[ 1.92295088 2.14601566] [ 2.02016082 0.3787193 ]

Help Variable Explorer Plots Files

Console 1/A X

\*c\* argument looks like a single numeric RGB or RGBA sequence, which should be avoided as value-mapping will have precedence in case its length matches with \*x\* & \*y\*. Please use the \*color\* keyword-argument or provide a 2D array with a single row if you intend to specify the same RGB or RGBA value for all points.

[[52 6]  
[ 3 19]]

Warning

Figures now render in the Plots pane by default. To make them also appear inline in the Console, uncheck "Mute Inline Plotting" under the Plots pane options menu.

In [5]:

Python Console History

LSP Python: ready conda: base (Python 3.9.13) Line 60, Col 13 UTF-8 CRLF RW Mem 64%



## Visualising the Test set results

Spyder (Python 3.9)

```

61 plt.show()
62
63 # 10: Visualising the Test set results
64
65 from matplotlib.colors import ListedColormap
66
67 X_set, y_set = X_test, y_test
68
69 X1, X2 = np.meshgrid(np.arange(start = X_set[:, 0].min() - 1, stop = X_set[:, 0].max() + 1, step = 0.01),
70                      np.arange(start = X_set[:, 1].min() - 1, stop = X_set[:, 1].max() + 1, step = 0.01))
71
72 plt.contourf(X1, X2, classifier.predict(np.array([X1.ravel(), X2.ravel()]).T).reshape(X1.shape),
73             alpha = 0.75, cmap = ListedColormap(('red', 'green')))
74
75 plt.xlim(X1.min(), X1.max())
76 plt.ylim(X2.min(), X2.max())
77
78 for i, j in enumerate(np.unique(y_set)):
79     plt.scatter(X_set[y_set == j, 0], X_set[y_set == j, 1], alpha=0.75,
80               c = ListedColormap(('red', 'green'))(i), label = j)
81
82 plt.title('Logistic Regression (Test set)')
83 plt.xlabel('Age')
84 plt.ylabel('Estimated Salary')
85
86 plt.legend()
87
88 plt.show()
89
90
91
92
93
94
95
96
97
98
99
100

```

Name	Type	Size	Value
classifier	linear_model._logistic.LogisticRegression	1	LogisticRegression object of s...
cm	Array of int64	(2, 2)	[[52 6] [ 3 19]]
i	int	1	1
j	int64	1	1
mydataset	DataFrame	(400, 5)	Column names: User ID, Gender, Age, EstimatedSalary, Purchased
sc	preprocessing._data.StandardScaler	1	StandardScaler object of sklea...
X	Array of int64	(400, 2)	[[ 19 19000] [ 35 20000]]
X1	Array of float64	(621, 586)	[[ -2.59882175 -2.58882175 -2.5... 3 ...
X2	Array of float64	(621, 586)	[[ -2.68064126 -2.68064126 -2.6... -2 ...
X_set	Array of float64	(80, 2)	[[ -0.49618606 0.56021375] [ 0.2389044 -0.59133674]]

Console I/O X

```

intend to specify the same RGB or RGBA value for all points.
*c* argument looks like a single numeric RGB or RGBA sequence, which should be avoided
as value-mapping will have precedence in case its length matches with *x* & *y*.
Please use the *color* keyword-argument or provide a 2D array with a single row if you
intend to specify the same RGB or RGBA value for all points.
[[52 6]
 [ 3 19]]
*c* argument looks like a single numeric RGB or RGBA sequence, which should be avoided
as value-mapping will have precedence in case its length matches with *x* & *y*.
Please use the *color* keyword-argument or provide a 2D array with a single row if you
intend to specify the same RGB or RGBA value for all points.
*c* argument looks like a single numeric RGB or RGBA sequence, which should be avoided
as value-mapping will have precedence in case its length matches with *x* & *y*.
Please use the *color* keyword-argument or provide a 2D array with a single row if you
intend to specify the same RGB or RGBA value for all points.
In [6]:

```

Python Console History

LSP Python: ready conda: base (Python 3.9.13) Line 63, Col 6 UTF-8 CRLF RW Mem 65%

